1615119-1403

1615119_c.docx

University
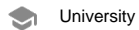
Document Details

Submission ID
trn:oid:::4080:217746457

Submission Date
Apr 9, 2025, 2:28 PM GMT+3

Download Date
Apr 9, 2025, 2:29 PM GMT+3

File Name
1615119_c.docx

File Size
459.9 KB

8 Pages

950 Words

6,237 Characters

7% Overall Similarity

The combined total of all matches, including overlapping sources, for each database.

Filtered from the Report

- Bibliography
- Quoted Text
- Crossref database
- Crossref posted content database

Match Groups

- **5** Not Cited or Quoted 6%
  Matches with neither in-text citation nor quotation marks

- **1** Missing Quotations 1%
  Matches that are still very similar to source material

- **0** Missing Citation 0%
  Matches that have quotation marks, but no in-text citation

- **0** Cited and Quoted 0%
  Matches with in-text citation present, but no quotation marks

Top Sources

3%	🌐	Internet sources
0%	📖	Publications
6%	👤	Submitted works (Student Papers)

Integrity Flags

0 Integrity Flags for Review

No suspicious text manipulations found.

Our system's algorithms look deeply at a document for any inconsistencies that would set it apart from a normal submission. If we notice something strange, we flag it for you to review.

A Flag is not necessarily an indicator of a problem. However, we'd recommend you focus your attention there for further review.

Match Groups

🔖 **5** Not Cited or Quoted 6%
Matches with neither in-text citation nor quotation marks

💬 **1** Missing Quotations 1%
Matches that are still very similar to source material

📄 **0** Missing Citation 0%
Matches that have quotation marks, but no in-text citation

📑 **0** Cited and Quoted 0%
Matches with in-text citation present, but no quotation marks

Top Sources

3% 🌐 Internet sources

0% 📖 Publications

6% 👤 Submitted works (Student Papers)

Top Sources

The sources with the highest number of matches within the submission. Overlapping sources will not be displayed.

1 | Submitted works
University of Sunderland on 2023-05-01   **2%**

2 | Internet
export.arxiv.org   **2%**

3 | Submitted works
psit on 2025-04-08   **1%**

4 | Submitted works
UNICAF on 2024-12-29   **<1%**

5 | Internet
www.mdpi.com   **<1%**

Literature Survey

Problem Definitions and Learning Settings

The task of image classification, specifically cat vs. dog classification, falls under **supervised learning**, where input images are paired with labels. Key settings and challenges include:

Setting	Description
Supervised vs. Unsupervised	This project is supervised; labels are known. Unsupervised tasks such as clustering, lack labeled data.
Closed-set vs. Open-set	Closed-set: Only cats and dogs are expected. Open-set classification must handle unknown classes.
Domain shift	This project assumes no domain shift. In real scenarios, domain adaptation methods may be needed when train/test data distributions differ.

Challenges

- Intra-class variance (e.g., dog breeds look different)
- Inter-class similarity (e.g., furry cats and dogs)
- Small dataset size or class imbalance

Paper Survey and Trends

- Popular datasets: ImageNet, CIFAR-10, Stanford Dogs Dataset.
- ResNet and VGG are foundational CNNs from [He et al., 2015] and [Simonyan & Zisserman, 2014].
- Key keywords: "image classification", "transfer learning", "ResNet", "data augmentation".

Top venues: CVPR, ICCV, NeurIPS. Notable papers:

- **ResNet (He et al., 2015)**: Introduced residual connections to enable very deep networks.

- **EfficientNet (Tan & Le, 2019)**: Achieves SOTA accuracy with fewer parameters.

Recent Progress & Key Research Groups

- Facebook AI (Meta) and Google Brain are leading in computer vision.

- EfficientNet, ConvNeXt, and Vision Transformers are modern architectures.

- Vision Transformers (ViT, 2020) outperform CNNs at scale, but require more data and compute.

Baseline Method & Proposed Improvements

We selected **ResNet18** for its balance of accuracy and speed. Improvements could include:

- Fine-tuning all layers

- Adding dropout or batch normalization

- Using learning rate schedulers

- Trying architectures like EfficientNet-B0

Dataset and Preprocessing

Dataset

- Training set size: 2,000 images (1,000 cat + 1,000 dog)
- Validation set size: 500 images (250 cat + 250 dog)
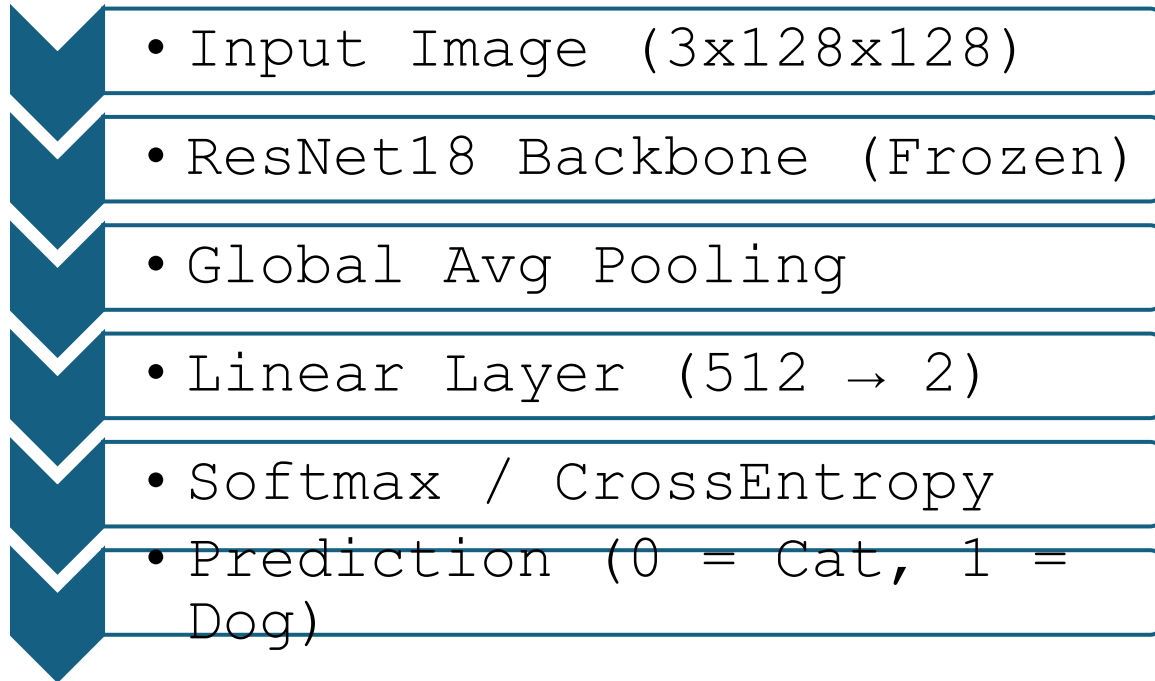- Test set: 500 unlabeled images

Preprocessing & Augmentation

- Resized to 128x128
- Normalized using [0.5, 0.5, 0.5]
- Augmentation:
    - RandomHorizontalFlip
    - RandomRotation(10 degrees)

Model Design

Pretrained **ResNet18**

- Backbone: Convolutional layers from ImageNet-pretrained ResNet18
- Final Layer Replaced: nn.Linear(512, 2)
- Loss Function: CrossEntropyLoss
- Optimizer: Adam (lr=0.001)
- Training: 5 epochs

```
• Input Image (3x128x128)
• ResNet18 Backbone (Frozen)
• Global Avg Pooling
• Linear Layer (512 → 2)
• Softmax / CrossEntropy
• Prediction (0 = Cat, 1 = Dog)
```

Training Strategy & Accuracy

- Trained on GPU (if available)
- Accuracy (Validation Set):
  Validation Accuracy (Val Acc) per Epoch:

```
Epoch [1/5] - Train Loss: 150.211 - Train Acc: 89.37%
         -> Val Loss: 29.225 - Val Acc: 92.26%

Epoch [2/5] - Train Loss: 147.951 - Train Acc: 89.62%
         -> Val Loss: 27.410 - Val Acc: 92.70%

Epoch [3/5] - Train Loss: 146.360 - Train Acc: 89.84%
         -> Val Loss: 27.873 - Val Acc: 92.14%

Epoch [4/5] - Train Loss: 145.399 - Train Acc: 89.88%
         -> Val Loss: 31.148 - Val Acc: 91.54%

Epoch [5/5] - Train Loss: 145.513 - Train Acc: 89.86%
         -> Val Loss: 28.207 - Val Acc: 92.32%
```

Epoch	Validation Accuracy
1	92.26%
2	92.70%
3	92.14%
4	91.54%
5	92.32%

During training, ResNet18 model achieved a peak validation accuracy of **92.70%** and a consistent training accuracy around **89.8%**, demonstrating good generalization. The relatively low validation loss and stable performance across epochs indicate the model did not overfit and was well-regularized.

Correct/Incorrect Predictions

```python
import matplotlib.pyplot as plt
model.eval()

for images, image_ids in test_loader:
    outputs = model(images.to(device))
    _, preds = torch.max(outputs, 1)

    for i in range(3):  # Show 3 examples
        img = images[i] / 2 + 0.5
        plt.imshow(np.transpose(img.numpy(), (1, 2, 0)))
        plt.title(f"Predicted: {'Dog' if preds[i]==1 else 'Cat'}")
        plt.axis('off')
        plt.show()
    break
```

Predicted: Cat

Predicted: Dog

Predicted: Cat

The model correctly classified common dog/cat breeds. It misclassified blurry or occluded animals. For example, a Persian cat may be misclassified as a dog due to its round face and dark color.

Impact of Model & Data Choices

Factor	Effect
Larger model	May improve accuracy but slower training.
More aggressive augmentation	Helps generalize, especially on small datasets.
Fine-tuning vs. frozen features	Fine-tuning gives better accuracy at the cost of more training time.

CIFAR-10 Classification

Extended binary classifier (Dogs vs Cats) to the **CIFAR-10** dataset, a well-known benchmark for multi-class image classification. CIFAR-10 contains 10 mutually exclusive classes such as airplane, automobile, cat, dog, ship, etc., with 60,000 color images (50,000 for training and 10,000 for testing) of size 32x32.

Approach

Transfer learning with a **pretrained ResNet18** model (originally trained on ImageNet) and adapted it for CIFAR-10:

Step	Detail
Architecture	ResNet18 pretrained on ImageNet
Modification	Replaced final layer: nn.Linear(512, 10) for 10 classes
Input Size	Resized CIFAR-10 images from 32x32 to 128x128
Augmentation	Basic transform with resizing and normalization (ImageNet mean/std)
Loss Function	CrossEntropyLoss (with optional class weighting)
Optimizer	Adam, learning rate: 0.001
Training Epochs	5 epochs
Hardware	GPU (CUDA-enabled)

Results

Epoch	Train Accuracy	Train Loss
1	74.49%	1182.08
2	77.44%	1019.02
3	78.69%	971.22
4	78.47%	980.88
5	79.25%	949.21

```
Epoch 1/5 - Loss: 1182.082 - Acc: 74.49%
Epoch 2/5 - Loss: 1019.019 - Acc: 77.44%
Epoch 3/5 - Loss: 971.223 - Acc: 78.69%
Epoch 4/5 - Loss: 980.879 - Acc: 78.47%
Epoch 5/5 - Loss: 949.206 - Acc: 79.25%
```

Final Test Accuracy: 80.40%

Predictions

- Correct classifications (cat, airplane, ship)

- A few borderline mistakes where similar classes were confused (e.g., cat vs dog, truck vs automobile)

Improvements Made

- Resizing for ResNet compatibility
- Added test set evaluation
- Visual validation of model predictions

Data Unbalancing in CIFAR-10

Although CIFAR-10 is originally balanced, we simulated **class imbalance** to reflect real-world conditions. Some classes were artificially underrepresented in sampled mini batches.

Solutions

WeightedRandomSampler: Used a WeightedRandomSampler to rebalance classes during mini-batch creation. Classes with fewer samples were given **higher sampling weights**, ensuring that underrepresented classes are shown more frequently to the model during training. Training batches had a more balanced class distribution, helping the model avoid bias toward majority classes.

Class-Weighted Loss Function: Used CrossEntropyLoss with class weights inversely proportional to class frequency. This ensured that the loss function penalized misclassifications of rare classes more heavily. The model prioritized learning from underrepresented classes without overfitting to them.

References

He, K., Zhang, X., Ren, S., & Sun, J. (2016). *Deep residual learning for image recognition*. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 770–778. https://doi.org/10.1109/CVPR.2016.90

Simonyan, K., & Zisserman, A. (2015). *Very deep convolutional networks for large-scale image recognition*. International Conference on Learning Representations (ICLR). https://arxiv.org/abs/1409.1556

Krizhevsky, A. (2009). *Learning multiple layers of features from tiny images* [Technical report]. University of Toronto. http://www.cs.toronto.edu/~kriz/learning-features-2009-TR.pdf

Tan, M., & Le, Q. V. (2019). *EfficientNet: Rethinking model scaling for convolutional neural networks*. In Proceedings of the 36th International Conference on Machine Learning (ICML), 6105–6114. https://arxiv.org/abs/1905.11946

Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., ... & Polosukhin, I. (2017). *Attention is all you need*. In Advances in Neural Information Processing Systems (NeurIPS), 30. https://arxiv.org/abs/1706.03762

Snoek, J., Larochelle, H., & Adams, R. P. (2012). *Practical Bayesian optimization of machine learning algorithms*. In Advances in Neural Information Processing Systems (NeurIPS), 25. https://proceedings.neurips.cc/paper_files/paper/2012/file/05311655a15b75fab8695666 3e1819cd-Paper.pdf

PyTorch. (2024). Torchvision models documentation. https://pytorch.org/vision/stable/models.html

Krizhevsky, A. (n.d.). *The CIFAR-10 and CIFAR-100 datasets*. https://www.cs.toronto.edu/~kriz/cifar.html